

# Rajeev Atla

AI/ML Engineer and Data Scientist Building Scalable, Secure, Smart Systems

US Citizen | [732-209-3995](tel:732-209-3995) | [rajeev.atla@gmail.com](mailto:rajeev.atla@gmail.com) | [linkedin.com/in/rajeev-atla](https://linkedin.com/in/rajeev-atla) | [github.com/RajeevAtla](https://github.com/RajeevAtla) | [rajeevatla.com](https://rajeevatla.com)

## EDUCATION

**Princeton University (Exchange Student)/Rutgers University - School of Engineering** Sep 2025 — May 2026  
*Master of Science in Computer Engineering (Specialization in Machine Learning)* Princeton/New Brunswick, NJ

Coursework: Multimodal AI, LLMs, Reinforcement Learning, Machine Vision, AI for Precision Health, High Performance Computing

**Rutgers University - School of Engineering** Sep 2021 — May 2025  
*Bachelor of Science (Triple Major) in Computer Engineering, Computer Science, and Data Science* New Brunswick, NJ

Recipient of the Eleanor and Samuel Sneath Endowed Merit Scholarship for Engineering Students

Coursework: AI, ML, Distributed Deep Learning, Data Science, Robotics and Computer Vision, Info and Network Security

## SKILLS

- **Programming Languages**: Python, R, SQL, Java, C/C++/CUDA, JavaScript/TypeScript, Rust, Bash, Go
- **AI/ML**: NumPy, PyTorch, JAX, TensorFlow, Keras, Pandas, Scikit-Learn, OpenAI API, LangChain/LangGraph, OpenCV, DSPy, RAG, HuggingFace (Transformers, Tokenizers, Datasets, Diffusers), vLLM, pgvector, Pydantic, FastAPI, NLTK, spaCy
- **Visualization**: Jupyter/Marimo, Matplotlib, Seaborn, Plotly, Tableau, PowerBI, React, Svelte
- **Cloud & DevOps**: AWS, Microsoft Azure, OCI, GCP, GitHub Actions, Jenkins, Docker, Kubernetes, Slurm
- **Tools & Databases**: PySpark, Hadoop/Hive, MongoDB, PostgreSQL, Git, Linux, Flink, Kafka, Jira, PyTest, Codex, Claude Code

## CERTIFICATIONS

- **AWS**: [Certified Cloud Practitioner](#), [Certified Machine Learning Specialist](#), [Certified AI Practitioner](#)
- **Oracle (OCI)**: [AI Foundations Associate](#), [Generative AI Professional](#), [Data Science Professional](#), [Vector AI Search Professional](#)

## WORK EXPERIENCE

**Software Development Engineering Intern - AI/ML** Sep 2026 — Dec 2026  
Amazon Seattle, WA

**Software Engineering Intern - AI & Data** May 2026 — Aug 2026  
Visa Austin, TX

- Built an AI PySpark/SQL query tool over **10 PB** for **30+ finance analysts**, converting questions into star-schema queries
- Created an LLM-driven dashboard that **cut reporting time 60%** and generated on-demand financial reports

**AI Engineering Intern** May 2024 — Sep 2024  
Atlait Inc. Remote

- Engineered PyTorch diffusion models for real-time predictions, **optimizing latency by 96ms** and enabling faster decision-making
- Optimized Airflow-Hadoop data pipeline to **speed up analytics delivery by 13%** in a cross-functional Agile environment

## PROJECTS

**DocuMint** <https://bit.ly/DocuMint>

- Built a 5-agent LangGraph API doc-modernizer with Gradio, achieved **90%+ modernization coverage** on sample docs, **cut manual edit time by 50%** with a **4-tab UX**, hardened with **8 deterministic Pytest cases**
- Authored multi-agent system with structured, **lifting modernization accuracy by 35%** and **trimming LLM spend by 20%**

**dexMCP** <https://bit.ly/dexmcp>

- Engineered MCP server with **5+ reusable tools** and **5+ Pydantic models** for **20+ typed fields** and **100% of tool inputs**
- Built asynchronous DSPy + LangChain clients to auto-discover tools and execute multi-step agentic workflows

**SuperconGAN** <https://bit.ly/3z7JaqZ>

- Built a PyTorch-based GAN to create synthetic superconductivity data of various materials, enhancing generative AI applications
- Released Python package on PyPI, achieving over **80,000 downloads** and widespread adoption

**raceformer** <https://bit.ly/raceformer>

- Engineered “Real-to-Sim” CUDA pipeline processing **30GB of multimodal sensor data** using JAX-based transformer model
- Achieved a **95% safety pass rate** by fine-tuning RL policy, **outperforming standard models by 35%** in risk avoidance