Rajeev Atla

AI/ML Engineer & Data Scientist applying cutting-edge techniques to build secure, scalable, complex systems

US Citizen | 732-209-3995 | rajeev.atla@gmail.com | github.com/RajeevAtla | linkedin.com/in/rajeev-atla | rajeevatla.com

Education

Rutgers University - School of Engineering Sep 2025 — May 2026
Master of Science *in Computer Engineering (Specialization in Machine Learning)* *New Brunswick, NJ*

Coursework: Reinforcement Learning, Multimodal AI, High Performance/Distributed Computing, Machine Vision

Rutgers University - School of Engineering Sep 2021 — May 2025
Bachelor of Science *(Triple Major) in Computer Engineering, Computer Science, and Data Science* *New Brunswick, NJ*

Recipient of the Eleanor and Samuel Sneath Endowed Merit Scholarship for Engineering Students

Coursework: AI, Distributed Deep Learning, Data Science, Statistical Learning, Computer Vision, Information and Network Security

Skills

- **Programming Languages**: Python, R, SQL, Java, C/C++/CUDA, JavaScript/TypeScript, Rust, Bash
- **AI/ML**: NumPy, PyTorch, JAX, TensorFlow, Keras, Pandas, Scikit-Learn, OpenAI API, LangChain/LangGraph, OpenCV, DSPy, RAG, HuggingFace (Transformers, Tokenizers, Datasets, Diffusers), vLLM, pgvector, Pydantic, FastAPI, NLTK, spaCy
- **Data Visualization**: Matplotlib, Seaborn, Plotly, Tableau
- **Cloud & DevOps**: AWS, Microsoft Azure, OCI, GCP, GitHub Actions (CI/CD Pipeline), Docker, Kubernetes
- **Tools & Databases**: Jupyter, PySpark, Kafka, Git, Linux (Ubuntu), PostgreSQL, MongoDB, Jira, Codex, Claude Code

Certifications

- **AWS**: Certified Cloud Practitioner, Certified Machine Learning Specialist, Certified AI Practitioner
- **Oracle (OCI)**: AI Foundations Associate, Generative AI Professional, Data Science Professional, Vector AI Search Professional

Work Experience

AI Engineering Intern May 2024 — Sep 2024
Atlait Inc. *Remote*

- Developed a Python-SQL compression script for form data, **reducing storage costs by 7%** for enterprise clients
- **Accelerated mean response time by 96 milliseconds** by integrating PyTorch inference models into Kafka microservices
- Created a **> 1TB** RAG system, utilizing A/B testing to evaluate and optimize AI-powered search accuracy
- Optimized CI/CD pipeline to **speed up build times by 13%** in an Agile environment, ensuring efficient development cycles

Projects

raceformer https://bit.ly/raceformer

- Used JAX-based transformer model to pretrain on CMHT driving dataset, achieving < 0.01 masked reconstruction loss
- Used Gymnasium-based RL environment to finetune and achieve < 2 minutes per lap of track time

dexMCP https://bit.ly/dexmcp

- Engineered Model Context Protocol (MCP) server exposing **5+ reusable tools** and **5+ Pydantic models**
- Implemented parameter validation across **20+ typed fields** and **100% of tool inputs**
- Built asynchronous clients using DSPy and LangChain to auto-discover tools and execute multi-step requests

DocuMint https://bit.ly/DocuMint

- Built a 5-agent LangGraph + Gemini API doc-modernizer with Gradio, achieved **90%+ modernization coverage** on sample docs, **cut manual edit time by 50%** with a **4-tab UX**, hardened with **8 deterministic pytest cases** and network-safe skips
- Authored a modular multi-agent system with structured prompts and severity-prioritized research, **lifting modernization accuracy by 35%** and **trimming LLM API spend by 20%**

SuperconGAN https://bit.ly/3z7JaqZ

- Built a PyTorch-based GAN to create synthetic superconductivity data of various materials, enhancing generative AI applications
- Extracted and processed **80,000+ dataset entries** from the UCI ML Repository using Pandas efficiently
- Released Python package on PyPI, achieving over **80,000 downloads** and widespread adoption
- Authored LaTeX paper detailing experimental design and validation methodology for **500,000+ data points**