

Rajeev Atla

US Citizen | [732-209-3995](tel:732-209-3995) | rajeevatla101@gmail.com | github.com/RajeevAtla | linkedin.com/in/rajeev-atla | rajeevatla.com

EDUCATION

Rutgers University - School of Engineering

Sep 2025 — May 2026

Master of Science in Computer Engineering (Specialization in Machine Learning)

New Brunswick, NJ

Coursework: Reinforcement Learning, Multimodal AI, High Performance/Distributed Computing

Rutgers University - School of Engineering

Sep 2021 — May 2025

Bachelor of Science (Triple Major) in Computer Engineering, Computer Science, and Data Science

New Brunswick, NJ

Recipient of the Eleanor and Samuel Sneath Endowed Merit Scholarship for Engineering Students

Coursework: AI, Distributed Deep Learning, Data Science, Statistical Learning, Computer Vision

SKILLS

- **Programming Languages:** Python, R, SQL, Java, C/C++/CUDA, Rust, Bash
- **AI/ML Libraries:** NumPy, PyTorch, JAX, TensorFlow, Keras, Pandas, Scikit-Learn, NLTK, LangChain/LangGraph
- **Data Visualization:** Matplotlib, Seaborn, Plotly, Tableau
- **Cloud & DevOps:** AWS, Microsoft Azure, OCI (Oracle Cloud Infrastructure), GitHub Actions, Docker, Kubernetes
- **Tools & Databases:** Jupyter Notebooks, Apache Kafka, Git, Linux (Ubuntu), PostgreSQL, MongoDB, Jira

CERTIFICATIONS

- **AWS:** [Certified Cloud Practitioner](#), [Certified Machine Learning Specialist](#), [Certified AI Practitioner](#)
- **Oracle (OCI):** [AI Foundations Associate](#), [Generative AI Professional](#), [Data Science Professional](#), [Vector AI Search Professional](#)

WORK EXPERIENCE

Software Engineering Intern

May 2024 — September 2024

Atlait Inc.

Remote

- Developed a Python-SQL compression script for form data, **reducing storage costs by 7%** for enterprise clients
- Integrated PyTorch inference into Kafka-microservices architecture, **improving mean response time by 96 milliseconds**
- Updated codebase from ES5 to ES7 using HTML, CSS, and TypeScript, resulting in **23% faster mean page loads**
- Optimized CI/CD pipeline to **speed up build times by 13%** ensuring efficient development cycles

PROJECTS

dexMCP

<https://bit.ly/dexmcp>

- Engineered Model Context Protocol (MCP) server exposing **5+ reusable tools** and **5+ Pydantic models**
- Implemented parameter validation across **20+ typed fields** and **100% of tool inputs**
- Built asynchronous clients using DSPy and LangChain to auto-discover tools and execute multi-step requests

SuperconGAN

<https://bit.ly/3z7JaQZ>

- Built a PyTorch-based GAN to model superconductivity data of various materials, enhancing generative AI applications
- Extracted and processed **80,000+ dataset entries** from the UCI ML Repository using Pandas efficiently
- Released Python package on PyPI, achieving over **80,000+ downloads** and widespread adoption
- Authored a LaTeX paper on findings and future scope, **incorporating 500,000+ data points** effectively

Cityscape (2nd Overall at HackExeter 2021)

<https://bit.ly/3OZjJ07>

- Designed and implemented a city tour mobile app, resulting in **100+ vivid city tours** for users
- Wrote controllers and models for MongoDB using MongooseORM to store **30+ kB of geographic data** in NoSQL schema
- Built mobile user interface allowing users to search, review, rank, and explore **100+ tours** using Flutter/Dart
- Constructed REST API using Express.js and nodemon to **increase development velocity by 20%** with hot-reloading

DocuMint

<https://bit.ly/DocuMint>

- Built a 5-agent LangGraph + Gemini doc-modernizer with Gradio, achieved **90%+ modernization coverage** on sample docs, **cut manual edit time 50%** with a **4-tab UX**, hardened with **8 deterministic pytest cases** and network-safe skips
- Authored modular agents (fetcher, analyzer, researcher, generator, quality-checker) with structured prompts and severity-prioritized research, **lifting modernization accuracy by 35%** and **trimming LLM API spend by 20%** through top-issue capping, content truncation, and batching