

# Rajeev Atla

Hard-working and Motivated Builder

US Citizen | [732-209-3995](tel:732-209-3995) | [rajeev.atla@gmail.com](mailto:rajeev.atla@gmail.com) | [linkedin.com/in/rajeev-atla](https://www.linkedin.com/in/rajeev-atla) | [github.com/RajeevAtla](https://github.com/RajeevAtla) | [rajeevatla.com](https://rajeevatla.com)

## EDUCATION

Rutgers University - School of Engineering

Sep 2025 — Present

*Master of Science in Computer Engineering (Specialization in Machine Learning)*

*New Brunswick, NJ*

Coursework: Multimodal AI, LLMs, Reinforcement Learning, Machine Vision, High Performance/Distributed Computing

Rutgers University - School of Engineering

Sep 2021 — May 2025

*Bachelor of Science (Triple Major) in Computer Engineering, Computer Science, and Data Science*

*New Brunswick, NJ*

Recipient of the Eleanor and Samuel Sneath Endowed Merit Scholarship for Engineering Students

Coursework: AI, ML, Distributed Deep Learning, Data Science, Robotics and Computer Vision, Info and Network Security

## SKILLS

- **Programming Languages**: Python, R, SQL, Java, C/C++/CUDA, JavaScript/TypeScript, Rust, Bash
- **AI/ML**: NumPy, PyTorch, JAX, TensorFlow, Keras, Pandas, Scikit-Learn, OpenAI API, LangChain/LangGraph, OpenCV, DSPy, RAG, HuggingFace (Transformers, Tokenizers, Datasets, Diffusers), vLLM, pgvector, Pydantic, FastAPI, NLTK, spaCy
- **Visualization**: Matplotlib, Seaborn, Plotly, Tableau, PowerBI, React
- **Cloud & DevOps**: AWS, Microsoft Azure, OCI, GCP, GitHub Actions (CI/CD Pipeline), Docker, Kubernetes, Slurm
- **Tools & Databases**: Jupyter, PySpark, Hadoop/Hive, Git, Linux, PostgreSQL, MongoDB, Jira, ROS2, Codex, Claude Code

## CERTIFICATIONS

- **AWS**: [Certified Cloud Practitioner](#), [Certified Machine Learning Specialist](#), [Certified AI Practitioner](#)
- **Oracle (OCI)**: [AI Foundations Associate](#), [Generative AI Professional](#), [Data Science Professional](#), [Vector AI Search Professional](#)

## WORK EXPERIENCE

AI Engineering Intern

May 2024 — Sep 2024

Atlait Inc.

*Remote*

- Developed a Python-SQL compression script for form data, **reducing storage costs by 7%** for enterprise clients
- Engineered PyTorch inference models for real-time predictions, **optimizing latency by 96ms** and enabling faster decision-making
- Created a **> 1TB** RAG-PySpark system, utilizing A/B testing to optimize AI-powered search and recommendation accuracy
- Optimized Airflow-Hadoop data pipeline to **speed up analysis by 13%** in an Agile environment, speeding up development

## PROJECTS

DocuMint

<https://bit.ly/DocuMint>

- Built a 5-agent LangGraph + Gemini API doc-modernizer with Gradio, achieved **90%+ modernization coverage** on sample docs, **cut manual edit time by 50%** with a **4-tab UX**, hardened with **8 deterministic pytest cases** and network-safe skips
- Authored a modular multi-agent system with structured prompts and severity-prioritized research, **lifting modernization accuracy by 35%** and **trimming LLM API spend by 20%**

dexMCP

<https://bit.ly/dexmcp>

- Engineered Model Context Protocol (MCP) server exposing **5+ reusable tools** and **5+ Pydantic models**
- Implemented parameter validation across **20+ typed fields** and **100% of tool inputs**
- Built asynchronous clients using **DSPy** and **LangChain** to auto-discover tools and execute multi-step requests

SuperconGAN

<https://bit.ly/3z7JaqZ>

- Built a PyTorch-based GAN to create synthetic superconductivity data of various materials, enhancing generative AI applications
- Extracted and processed **80,000+ dataset entries** from the UCI ML Repository using Pandas efficiently
- Released Python package on PyPI, achieving over **80,000 downloads** and widespread adoption

raceformer

<https://bit.ly/raceformer>

- Engineered a high-fidelity “Real-to-Sim” validation pipeline processing **30GB of multimodal sensor data** (LiDAR, camera, radar) on 4x A100s, utilizing JAX-based vision-language model to generate ground truth scenarios for critical edge case simulation
- Achieved a **95% pass rate on safety metrics** by leveraging geometric priors to fine-tune RL policies, establishing clear performance baselines and **outperforming standard models by 35%** in neural path planning and risk avoidance